

Krishna Tripathi

+91 70160 50711 | kanchanlatakrishna@gmail.com | [LinkedIn](#) | [GitHub](#) | [LeetCode](#) | [Portfolio](#)

Education

Indian Institute of Information Technology, Nagpur

December 2021 to May 2025

Bachelor of Technology in Computer Science and Engineering

Nagpur, India

Technical Skills

Languages: C++, Go, Java, Python, JavaScript, TypeScript, SQL

Backend: Node.js, Express.js, FastAPI, REST APIs, Microservices

Frontend: React, Next.js **AI/LLM:** LangGraph, MCP, RAG, Langfuse

Databases: YugabyteDB, PostgreSQL, MongoDB, BigQuery, pgvector

Cloud/DevOps: AWS (Cognito, Lambda, S3), Google Cloud Run, Docker, Jenkins, GitHub Actions

Professional Experience

Yugabyte

Bengaluru, India

Software Engineer I

September 2025 to Present

Software Engineer Intern

September 2024 to August 2025

- **Engineered Meko knowledge base uploads across 3 services** using React, Go, Python, and presigned S3 POST requests. Enforced 3 storage edge checks for file size, type, and content spoofing while offloading large transfers from API servers.
- **Built Meko's memory retrieval observability and promotion flow** with top-k retrieval across memory and knowledge base context. Replaced full conversation replay, reducing input token usage by **35%** for marketing and lead generation workflows.
- **Delivered AWS Cognito authentication across 3 identity providers:** Google, LinkedIn, and GitHub. Supported profile completion and identity linking while hardening invite, logout, and federated account flows.
- **Contributed AI/RAG integrations across Java, Python, and C++:** YugabyteDB embedding storage and chat memory in LangChain4j, plus a PgDistRagRetriever and supporting tserver changes for [pg.dist.rag](#). Added HNSW and hybrid search, metadata filters, and comprehensive integration tests; published [docs](#) through a merged upstream [commit](#).
- **Developed a database migration validator for 2 health signals:** schema version and dirty state. Blocked invalid migrations before merge and surfaced actionable GitHub feedback.
- **Engineered xCluster disaster recovery services in Go and UI**, optimizing DR node selection from **450ms to 120ms** and exposing C++ replication lag metrics through Prometheus and UI monitoring.
- **Automated a Go analytics pipeline across 5 systems:** Scarf, GCS, BigQuery, MadKudu, and Salesforce. Ingested Scarf data through GCS into BigQuery for transformation, lead scoring, and CRM synchronization.
- **Expanded the YB Voyager experience in yugabyted UI across multiple open-source commits**, including the end-to-end workflow [implementation](#). Built assessment UX, stage status visibility, and progress analytics across assessment, schema, data, and cutover, simplifying migration tracking without CLI output.

Zluri

Bengaluru, India

Software Engineer Intern

August 2024 to September 2024

- **Resolved 13+ production bugs** in JavaScript/Node.js and MongoDB. Built QA dashboards and a testing framework, increasing coverage by **20%** and query performance by **30%**.

Projects

EternalAgni: Agnihotra Reminder App

July 2026 to Present

JavaScript, Capacitor, Android, PWA, Service Workers, Local Notifications

[SOURCE CODE](#) | [INTERNAL TESTING](#)

- **Built a custom NOAA timing engine** with **±1-second** accuracy. Generated sunrise and sunset schedules for **92 days (~3 months)** locally for offline use.
- **Delivered reliable reminders** with **2-to-60-minute** lead times. Added duplicate suppression, late delivery protection, and client-side scheduling.
- **Shipped releases across 2 platforms:** Android and PWA. Added GPS/IP fallbacks, service worker recovery, and automated APK/AAB builds.

YB Voyager Copilot: AI Powered Database Migration Assistant

January 2025 to Present

LangGraph, FastAPI, Next.js, Python, MCP, Docker, YugabyteDB, RAG

[DOCKER HUB](#) | [ARCHITECTURE](#) | [DEMO](#)

- **Architected a 3-agent LangGraph and MCP workflow** covering all 4 migration stages with grounded product documentation.
- **Built a 3-tier RAG backend and containerized 5 services** with FastAPI, pgvector, YugabyteDB, and Docker. Enabled persistent data across the workflow.
- **Automated 2 report workflows:** delivery and migration. Used GitHub Actions and Playwright to generate production sizing recommendations.

Achievements

- Received **Best Use of AI** recognition for YB Voyager Copilot at the YugabyteDB internal hackathon.
- Earned **Guardian status with a maximum rating of 2145** on LeetCode. Reached **4-star status with a 1944 rating** on CodeChef and **Specialist status with a maximum rating of 1484** on Codeforces.
- Secured **3rd place** in Algorithmia at Tantrafiesta 2023.